

Laboratory on processes

OS - Operating System

Simone ROSSI

October 24, 2016

Date Performed: 13/20 October 2016
Partners: Simone Rossi

1 Information on Operating Systems

1.1 History of Windows

According to the Wikipedia page of Windows NT, it is written in C, C++ and Assembly language. The spectrum of programming languages used is quite wide and is needed to explore different levels of abstraction with the right language. For example, probably boot operations and high performance applications (like “critical” drivers, such as memory management) were been written in assembly, while “normal” drivers (power management, audio, video, network, ...) and system libraries were been written in a higher level language like C. Finally, GUI may been written in C++.

1.2 Linux Kernel

- 1 `uname -a` prints system informations, such as the version of the operating system and the name of the machine in use. To print only the system's version the command to be run is `uname -v`, which in my case showed
121-Ubuntu SMP Wed Jan 20 10:50:42 UTC 2016
- 4 `/home/Local.Data` is a local directory on the PC, since is not part of the mounted filesystem
- 6 I didn't know the difference between “z” compression and “tar” compression. After a fast research, seems that tar compression is not a real compression: it is just an utility to collect many files into just one archive; on the contrary, with “z” the data is compressed with the LZW algorithm.
- 7 Yes, all the source file are prepared to be configure to target a specific platform. After the configuration, which I left with default parameters,

the kernel has been compiled. At the end, to use the new kernel, it has to be mount on a new root partition.

2 Analyzing processes

2 Here is reported the result of the command `vmstat -s`.

```
      0 swap ins
      0 swap outs
      0 pages swapped in
      0 pages swapped out
783503 total address trans. faults taken
      550 page ins
      0 page outs
      783 pages paged in
      0 pages paged out
296894 total reclaims
296894 reclaims from free list
      0 micro (hat) faults
783503 minor (as) faults
      550 major faults
167649 copy-on-write faults
290103 zero fill page faults
216060 pages examined by the clock daemon
      0 revolutions of the clock hand
      0 pages freed by the clock daemon
    1269 forks
      850 vforks
     2020 execs
31435117 cpu context switches
51201789 device interrupts
    976804 traps
10213769 system calls
3874932 total name lookups (cache hits 96%)
    11256 user   cpu
    79993 system cpu
30751051 idle   cpu
      0 stolen cpu
    1306 forks
      850 vforks
```

All the these informations are statistics that the OS keeps track since bootup: some of them are related to memory management (like pages in and out) and others to forks and system calls. Since bootup, the system has executed 1269 **forks** and 850 **vforks**.

4 The compilation ended with two errors:

```
fork01.c: In function  main  :
fork01.c:4:3: error: unknown type name  pid_t
```

```
fork01.c:19:1: error: expected      ;      before      }
```

These errors are solved including the library `<unistd.h>` and a semi-column in line 18.

5 To simply get rid the warning I added a `return 0;`.

6 There is a run time error, a segmentation fault. Looking at the strings of the executable file, this is the result:

```
d4 /usr/lib/ld.so.1
9b8 child
9be parent
9c5 process %s ret = %ld
b7e YZ]
bb8 []
c01 []
c18 &
cbd $ U
d64 [^]
```

The strings we are interested in start with offset 0009B8. Reading the ELF file, I found that the correspondent section is read-only and it has not the privilege to be written

```
[ 8] .SUNW_dynsymSORT L00S+ffffff1 08050918 000918
000038 04 A 3 0 4
[ 9] .SUNW_reloc REL 08050950 000950
000028 08 A 4 0 4
[10] .rel.plt REL 08050978 000978
000040 08 AI 4 16 4
[11] .rodata PROGBITS 080509b8 0009b8
000024 00 A 0 0 1
[12] .plt PROGBITS 080509dc 0009dc
000090 10 AX 0 0 4
[13] .text PROGBITS 08050a70 000a70
0002fc 00 AX 0 0 16
```

These lines, finally, are dump of the binary file in the section `.rodata`.

```
Contents of section .rodata:
80509b8 6368696c 64007061 72656e74 0070726f child.
parent.pro
80509c8 63657373 20257320 72657420 203d2025 cess %s
ret = %
80509d8 6c640a00 ld..
```

7 To make a more robust code, I simply added a control on the return value of the fork.

```

if(ret < 0){
    printf("Error in process creation");
    return -1;
}

```

8 These are a few lines of a long list of function and system calls:

CPU	ID	FUNCTION:NAME
0	81524	tsort_unsigned:entry
0	81523	tsort:entry
0	81461	set_environ:entry
0	81333	lookup_sym:entry
0	81353	elf_hash:entry
0	81332	_lookup_sym:entry
0	81330	core_lookup_sym:entry
0	81460	callable:entry
0	81354	elf_find_sym:entry
0	81755	strcmp:entry
0	81755	strcmp:entry
0	81755	strcmp:entry
0	81394	calloc:entry

Some of them are system calls (like `getpid`) and others are function calls (like `mutex.lock` and `mutex.unlock`). After that, I monitored the system call of my program:

```

forksys
fstatat64
ioctl
lwp_sigmask
nanosleep
rexit
schedctl
write

```

For instance, `forksys` refers to the `fork` function and `nanosleep` to sleep function. To run the provided script, I also wrote a little script to automate the retrieve of the PID

```

./rossi_fork &\
read pid <<< $( ps | grep rossi_fo | awk '{print $1;}')
;\
echo $pid;\
./dtracescript $(echo $pid) >> 2x08.txt

```

The result of this script is the following (just only few lines).

```

__forkx
__lwp_sigmask
__nanosleep
__schedctl

```

```

__syscall
__write
_calloc
_cleanup
_exithandle
_fflush_l_iops
_fflush_u
_findbuf
_getfp
_malloc
_ndoprnt
_postfork_parent_handler
_prefork_handler
_setbufend
_setorientation
_ti_critical
_xflsbuf
aplist_append
aplist_delete
aplist_insert
aplist_test

```

- 9 Using the same script and adding a couple of instruction in the source file, I can monitor the right moment in which the child is killed.

```

forkx
free
fstat64
fstatat64
ioctl
isatty
isseekable
kill
libc_fini

```

- 10 Using the provided script, this is the resut.

```

SIGKILL was sent to fork_palmiero by Christian Palmiero
SIGKILL was sent to rossi_fork by Simone Rossi
SIGKILL was sent to fork_palmiero by Christian Palmiero

```

- 11 The program I wrote is the following.

```

#include <stdio.h>
#include <unistd.h>
#include <signal.h>

#define MAX 3
#define SPL 10
int main(void) {

```

```

int i;
int j;

for (i = 0; i < MAX; i++)
{
    pid_t ret = fork();
    if (ret == 0)
    {
        if ( i == 0 || i == MAX - 1 )
        {
            for (j = 0; j < 2; j++)
            {
                pid_t ret2 = fork();
                if (ret2 == 0)
                {
                    printf("%d", j);
                    sleep(SPL);
                    exit(0);
                }
            }
            printf("%d/n", i);
            sleep(SPL);
            exit(0);
        }
    }
    sleep(200);
    printf("Parent");
    return(0);
}

```

Checking with ps and dtrace I verified that It correctly spawns three children.

CPU	ID	FUNCTION:NAME
0	80418	fork:entry
3	80418	fork:entry
1	80418	fork:entry

I also traced two function calls (printf and exit) and one system call (write)

CPU	ID	FUNCTION:NAME
0	80158	exit:entry

CPU	ID	FUNCTION:NAME
1	80159	printf:entry

CPU	ID	FUNCTION:NAME
0	80159	write:entry

- 12 To display information about the number and/or types of processors installed on the system, the command to run is `psrinfo -p -v`. For the machine `calibra`, the result is the following:

```
The physical processor has 2 virtual processors (0-1)
  x86 (GenuineIntel 206D7 family 6 model 45 step 7 clock
    1800 MHz)
    Intel(r) Xeon(r) CPU E5-2650L 0 @ 1.80GHz
The physical processor has 2 virtual processors (2-3)
  x86 (GenuineIntel 206D7 family 6 model 45 step 7 clock
    1800 MHz)
    Intel(r) Xeon(r) CPU E5-2650L 0 @ 1.80GHz
```

To map a process to a given set of virtual (or physical) processors the command to run is `pbind -b -c 0-1 'pid'`. In this case, the “pid” process is mapped to virtual processors 0 and 1 (physical processor 0).

3 Communication between processes

- 2 In the original program, only the “sender” creates the file while the “receiver” opens it only in read mode. This will bring to a condition race: the order of execution of the two processes is not subjected to any constrain: the scheduler can first schedule the sender and then the receiver or in the opposite order. In the first case, everything work properly while in the second case the receiver tries to open a file which doesn’t exist yet. The solution could be introduce a synchronization between processes (likely using semaphores) or create a delay in the receiver before it starts to work on the queue (using for example a `sleep` system call).

```
int main(void) {
    int pid;
    // Create two processes: one for sending data, and
    // another one for receiving data
    pid = fork();
    if (pid != 0) { /* Parent */
        sleep(1);
        receiver();
        // delete the message queue
        mq_unlink(MQ_NAME);

    } else { /* Child */
        sender();
    }
    return(0);
}
```

Finally, this is the `dtrace` result run on this program (the complete dump is provided as additional material).

```

./mqtest
I am in sender
0 ?--->
1 ?--->
2 ?--->
3 ?--->
4 ?--->
5 ?--->
6 ?--->
7 ?--->
8 ?--->
9 ?--->
10 ?--->
11 ?--->
12 ?--->
13 ?--->
14 ?--->
I am in receiver
---> 0
---> 1
---> 2
---> 3
---> 4
---> 5
---> 6
---> 7
---> 8
---> 9
---> 10
---> 11
---> 12
---> 13
---> 14
CPU      ID      FUNCTION:NAME
  2   85273      tsort_unsigned:entry
  2   85272          tsort:entry
  2   85210      set_environ:entry
  2   85082      lookup_sym:entry
  2   85102          elf_hash:entry
  2   85081      _lookup_sym:entry

```

- 3 For this exercise, there are three main actors: a sender (son11), a forwarder (son2) and a receiver (son12).

This is source code of the sender:

```

int n;
mqd_t mqfd;

//printf("I am in sender\n");
init_queue (&mqfd, name, O_CREAT | O_WRONLY);

```



```

for (n = 0; n < MAX_MSG; n++) {
    printf ("%d?--->\n", n);
    put_integer_in_mq (mqfd, n);
}
put_integer_in_mq (mqfd, OVER);

/* close queue */
mq_close (mqfd);

```

This is source code of the forwarder:

```

int data_from_s11;
mqd_t mq11_2, mq2_12;

init_queue (&mq2_12, "/mq2_12", O_CREAT | O_WRONLY);
init_queue (&mq11_2, "/mq11_2", O_CREAT | O_RDONLY);

while (1) {
    data_from_s11 = get_integer_from_mq (mq11_2);
    put_integer_in_mq(mq2_12, data_from_s11);

    if (data_from_s11 == OVER)
        break;

    printf("--->_%d--->\n", data_from_s11);
}

```

This is source code of the receiver:

```

int d;
mqd_t mqfd;

// printf("I am in receiver\n");
init_queue (&mqfd, name, O_CREAT | O_RDONLY);
while (1) {
    d = get_integer_from_mq (mqfd);
    if (d == OVER)
        break;
    printf ("--->_%d\n", d);
}
/* Close queue */
mq_close(mqfd);

```

This is the source code needed to spawn the correct process tree:

```

int main(void) {
    pid_t ret1, ret2;

    ret1 = fork();
    if (ret1 < 0)
        perror("Error in process creation: SON1 cannot be
            created");
}

```

```

else if (ret1 == 0)
    son1();
else {
    ret2 = fork();
    if (ret2 < 0)
        perror("Error_in_process_creation:_SON2_cannot_
            be_created");
    else if (ret2 == 0)
        son2();
    else {
        int i;
        for (i = 0; i < 2; i++)
            wait(NULL);
    }
}

return(0);
}

void son1(void) {
    pid_t ret11, ret12;

    ret11 = fork();
    if (ret11 < 0)
        perror("Error_in_process_creation:_SON11_cannot_be_
            created");
    else if (ret11 == 0)
        son11();
    else {
        ret12 = fork();
        if (ret12 < 0)
            perror("Error_in_process_creation:_SON12_cannot_
                be_created");
        else if (ret12 == 0)
            son12();
        else {
            int i;
            for (i = 0; i < 2; i++)
                wait(NULL);
            exit(0);
        }
    }
}

```

This is the terminal dump of a run of this program:

```

0 ?--->
1 ?--->
2 ?--->
3 ?--->
4 ?--->
5 ?--->
6 ?--->

```

```

7 ?--->
8 ?--->
9 ?--->
10 ?--->
11 ?--->
12 ?--->
13 ?--->
14 ?--->
----> 0 ---->
----> 1 ---->
----> 2 ---->
----> 3 ---->
----> 4 ---->
----> 5 ---->
----> 6 ---->
----> 7 ---->
----> 8 ---->
----> 9 ---->
----> 10 ---->
----> 11 ---->
----> 12 ---->
----> 13 ---->
----> 14 ---->
----> 0
----> 1
----> 2
----> 3
----> 4
----> 5
----> 6
----> 7
----> 8
----> 9
----> 10
----> 11
----> 12
----> 13
----> 14

```

Bonus For the bonus question, the same forwarding operation is now performed also by the son1 and this is the source code to do that:

```

int data_from_s11;
mqd_t mq11_1, mq1_12;

init_queue (&mq1_12, "/mq2_12", O_CREAT |
O_WRONLY);
init_queue (&mq11_1, "/mq11_2", O_CREAT |
O_RDONLY);

```

```

while (1) {
    data_from_s11 = get_integer_from_mq (mq11_1);
    put_integer_in_mq(mq1_12, data_from_s11);

    if (data_from_s11 == OVER)
        break;

    printf("---->s1:_%d_---->\n", data_from_s11);
}

int i;
for (i = 0; i < 2; i++)
    wait(NULL);
exit(0);

```

The only different with respect to the previous program is the inclusion of two `OVER` number from the sender. This is need to guarantee that both `son1` and `son2` terminate their executions.

```

int n;
mqd_t mqfd;

//printf("I am in sender\n");
init_queue (&mqfd, name, O_CREAT | O_WRONLY);
for (n = 0; n < MAX_MSG; n++) {
    printf ("%d_?---->\n", n);
    put_integer_in_mq (mqfd, n);
}
put_integer_in_mq (mqfd, OVER);
put_integer_in_mq (mqfd, OVER);
/* close queue */
mq_close (mqfd);
return ;

```

Finally, this is the terminal dump of this last program.

```

0 ?---->
1 ?---->
2 ?---->
3 ?---->
4 ?---->
5 ?---->
6 ?---->
7 ?---->
8 ?---->
9 ?---->
10 ?---->
11 ?---->
12 ?---->
---->s1: 1 ---->

```

```

---->s2: 0 ---->
---->s1: 2 ---->
---->s2: 3 ---->
13 ?---->
---->s1: 4 ---->
---->s2: 5 ---->
---->s2: 7 ---->
14 ?---->
---->s1: 6 ---->
---->s2: 8 ---->
---->s1: 9 ---->
---->s2: 10 ---->
----> 0
----> 1
----> 2
----> 3
----> 4
----> 5
---->s1: 11 ---->
----> 6
----> 7
----> 8
----> 9
---->s2: 12 ---->
---->s2: 14 ---->
---->s1: 13 ---->
----> 10
----> 11
----> 12
----> 13
----> 14

```